Adaptive Aggregation of Recommender Systems

Olav Bjørkøy

Norwegian University of Science and Technology

Department of Computer and Information Science

India-Norway Workshop on

Web Concepts and Technologies

October 3rd, 2011

Terminology

term	description
u	a user
i	an item (article, website, movie, email...)
r	the rating (relevance, utility, other domain term)
m	a method/algorithm for predicting ratings.
$p(m, u, i)$	rating prediction from method m for (u, i)

2006: The Netflix Challenge

- USD 1MM prize for a 10% accuracy improvement.
- Breakthrough: Combining methods from many teams.
- Team BellKor finally achieved a 10.06% improvement by combining **107** different recommender algorithms.

Today: Web Search

"Today we use more than 200 signals, including PageRank, to order websites, and we update these algorithms on a weekly basis."
— Google

(google.com/corporate/tech.html)

"We use over 1,000 different signals and features in our ranking algorithm."
— Microsoft Bing

(bing.com/community/site_blogs/b/search/archive/2011/02/01/
thoughts-on-search-quality.aspx)

Why multiple algorithms?

- ▶ The unreasonable effectiveness of data
- ▶ Capture more predictive aspects of existing data.
- ▶ Specialized predictors for subsets of data.

"Quite frequently we have found that the more accurate predictors are less useful within the full blend."
— Bell, R., Koren, Y., and Volinsky, C. (2007) (Netflix)

The Problem: Latent Subjectivity

$$\hat{r}_{u,i} = \sum_{m \in M} w_m \times p_r(m, u, i) \tag{1}$$

- ▶ Generalized optimal weights.
- ▶ Treats all users and items the same.
- ▶ Varying accuracy across users and items.
- ▶ Methods are chosen by the system, not the users or items.

Systems that insist on being adaptive in a certain way are not really adaptive at all.

Adaptive Recommenders

- Combining predictors by estimated contextual accuracy.

$$\hat{r}_{u,i} = \sum_{m \in M} p_w(m, u, i) \times p_r(m, u, i) \qquad (2)$$

- p_r: predicted rating from method m for (u, i).
- p_w: predicted optimal weight for method m for (u, i).
- We can use recommender systems for **both** p_r and p_w.
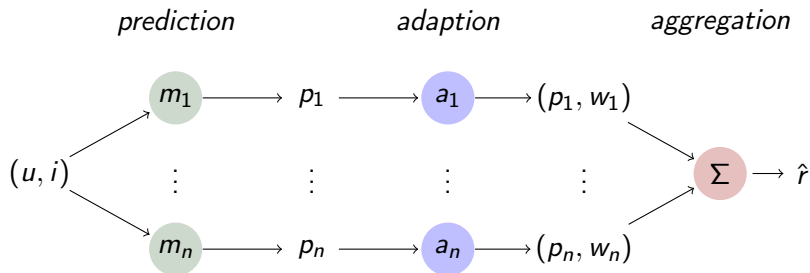
Adaptive Recommenders

Figure: Layers of recommenders.

Training phase

1. Split data into two sets: (d_m, d_e).
2. Use d_m to train the rating predictors.
3. Create *error matrices* for the rating predictors and d_e.
4. The error matrices now hold known errors for some (m, u, i).
5. Train error predictors from the error matrices.

$$\forall(u, i, r) \in (d_e - d_m) : E(m)_{u,i} = |r - p(m, u, i)| \qquad (3)$$

Training phase

$$Rating_{u,i} = \begin{pmatrix} r_{1,1} & r_{1,2} & \cdots & r_{1,i} \\ r_{2,1} & r_{2,2} & \cdots & r_{2,i} \\ \vdots & \vdots & \ddots & \vdots \\ r_{u,1} & r_{u,2} & \cdots & r_{u,i} \end{pmatrix}$$

$$Error(m)_{u,i} = \begin{pmatrix} e_{1,1} & e_{1,2} & \cdots & e_{1,i} \\ e_{2,1} & e_{2,2} & \cdots & e_{2,i} \\ \vdots & \vdots & \ddots & \vdots \\ e_{u,1} & e_{u,2} & \cdots & e_{u,i} \end{pmatrix}$$

- $\mathrm{train}(m_1, R) \rightarrow$ rating predictor for m_1
- $\mathrm{train}(m_2, E(m_1)) \rightarrow$ error predictor for m_1

Prediction phase

1. Predict ratings $\hat{r}_{(u,i,m)}$.
2. Predict errors $\hat{e}_{(u,i,m)}$.
3. Create adaptive weights by inversing the normalized errors.
4. Sum the weighted predictions to get the final $\hat{r}$.

Prediction phase

$$\hat{r}_{u,i} = \sum_{(m_e, m_r) \in M} (1 - \frac{p(m_e, u, i)}{error(u, i)}) \times p(m_r, u, i) \quad (4)$$

$$error(u, i) = \sum_{m_e \in M} p(m_e, u, i) \quad (5)$$

Experiment

- RMSE values for basic recommenders, simple aggregations and adaptive aggregation.
- Used the Movielens movie rating dataset.
- See paper for more details.

$$\mathrm{RMSE}(\hat{R}, R) = \sqrt{\frac{\sum_{i=1}^{n}(\hat{R}_i - R_i)^2}{n}} \qquad (6)$$

Results

(a) RMSE values for the five disjoint subsets:

	method	d_1	d_2	d_3	d_4	d_5
S	svd1	1.2389	1.1260	1.1327	1.1045	1.1184
S	svd2	1.2630	1.1416	1.1260	1.1458	1.1260
S	svd3	1.0061	0.9825	0.9830	0.9815	0.9797
S	svd4	1.0040	0.9830	0.9849	0.9850	0.9798
S	slope_one	1.1919	1.0540	1.0476	1.0454	1.0393
S	item_avg	1.0713	0.9692	0.9662	0.9683	0.9725
S	baseline	1.0698	0.9557	0.9527	0.9415	0.9492
S	cosine	1.1101	0.9463	0.9412	0.9413	0.9382
S	knn	1.4850	1.1435	1.1872	1.2156	1.2022
A	median	0.9869	0.8886	0.8857	0.8857	0.8855
A	average	0.9900	0.8536	0.8525	0.8525	0.8519
A	adaptive	**0.9324**	**0.8015**	**0.7993**	**0.8238**	**0.8192**

(b) Statistics for the methods:

|	method	min	max	mean	σ	Δ
S	knn	1.1435	1.4850	1.2467	0.3487	-
S	svd2	1.1260	1.2630	1.1605	0.2277	6.9%
S	svd1	1.1045	1.2389	1.1441	0.2197	1.4%
S	slope_one	1.0393	1.1919	1.0756	0.2415	5.9%
S	item_avg	0.9662	1.0713	0.9895	0.2023	8.0%
S	svd4	0.9798	1.0040	0.9873	**0.0924**	2.2%
S	svd3	0.9797	1.0061	0.9865	0.0991	0.1%
S	cosine	0.9382	1.1101	0.9754	0.2595	1.1%
S	baseline	0.9415	1.0698	0.9738	0.2196	1.6%
A	median	0.8855	0.9865	0.9065	0.2005	6.9%
A	average	0.8519	0.9900	0.8801	0.2344	2.9%
A	adaptive	**0.7993**	**0.9324**	**0.8352**	0.2225	5.1%

Results

Figure: Average RMSE plot.

Results


Figure: RMSE Standard deviation caused by dataset d_1.

Limitations

- Lots of added complexity for fairly unknown improvement.
- Only tested on a few datasets, no real world situations.
- Only compared to simple aggregation methods.
- Neither the aggregators nor the basic recommenders were heavily optimized to the domain of the dataset.

Adaptive Recommenders

- Combine disjoint algorithms
- Weigh recommenders by predicted accuracy.
- Accuracy predictions are contextually dependent on (u, i, m).
- *Any* applicable recommender becomes a worthy addition.

See paper for more references and results.